

IP Misaddress

Greenbook Club has come into possession of an old IPv4 networking system made by Swansea Uni researchers in the 80s (don't ask us how, John Tucker has his ways). The only problem is, some bright spark decided to break the spec, so nothing stored looks like IP addresses!

For example, Swansea's IP is normally 137.44.1.20, however this system has it stored as "zp.Zt.B.U"! That's right, it's a mnemonic IP system! Completely out of spec, but Tucker wants us to convert these old "Swansea IP" addresses into normal IPv4. Handily, there's an old comment left in a user manual.

```
/*                                                                 *\  
*                               Swansea IPv4 - ** SIP Address **   *  
*-----*  
* This system, unique to Swansea, is an extension on the ARPANET IPv4 way of *  
* addressing networks. It provides a way to use plaintext addresses for use *  
* on C based ASCII operating systems (ie. UNIX System V, because this is all *  
* we can afford). Thanks to its incomprehensible nature and hence similarity *  
* to Welsh, this extension proved popular with faculty and students. Well at *  
* least the ones outside the Welsh Studies department. To convert from IPv4: *  
* *  
* For each byte, starting with the leading byte, produce printable ASCII in *  
* the range 97-122 (a-z). If the leading byte is less than 97, then we use *  
* an uppercase display in the range 65-90 (A-Z), however we always interpret *  
* them as if in lowercase, but with the knowledge that we have offset this. *  
* This means a byte < 97 is first increased by 97 (upcast to avoid rollover) *  
* but is displayed in uppercase (A-Z) so that we can recognise this state. *  
* For a byte > 122, we now have two phases. We would like to print out chars *  
* that "sum" to the byte (with following char as 0-24 for a-z), however this *  
* would mean that we get a large string of 'z' characters in many cases. For *  
* this reason, we first convert these 'z' characters into decimal digits in *  
* the range 48-53 (0-5) such that, if there is more than one 'z', we produce *  
* the decimal representation of how many were required. Hence, a byte=146 is *  
* the lowest byte with this behaviour, and so contributes the char '2'. If *  
* the byte is lower than 97, we consider the first Z offset, so for a byte *  
* = 49, which is expanded up to 146 for a-z representation, however we force *  
* the first character to be represented as a 'Z' since this is needed so we *  
* can disambiguate the process. Following this would be a normal 'z', as no *  
* further characters are required. Once we reach the byte = 74 and above, we *  
* resume this compression scheme and produce the char '2' or greater. Due to *  
* the limitations of a single byte, this should not exceed '5'. Hence 255 is *  
* converted to "5i". This is repeated for each byte, which now do not need a *  
* '.' to be included between them, since this is unambiguous, however it can *  
* be included to produce "phonetic groupings" or such to ease readability. *  
* *  
* Example: 137.44.1.20 = "zp.Zt.B.U", 255.0.0.1 = "5iAAB"
```

```

*
* To convert back to IPv4, we simply invert this process.
* This is left as an exercise for the reader / unwitting CS-115 student...
*
*-----*
*      COPYRIGHT - 1985- under the Swansea Fictitious Copyright License
*///                                     \\\*

```

As an extension, some IP addresses in our modern table aren't in the normal "quad" format (like 137.44.1.20), but instead are other compatible (in-spec) formats! So, in order to make matching easier (because who on earth canonises anything) Tucker wants you to produce each address in a range of different formats.

(Inspired by this fun trick & `inet_aton(3)` shenanigans.)